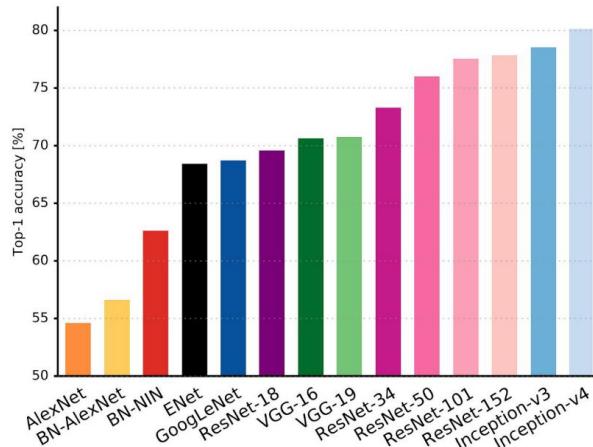


# **Neural Architecture Search**

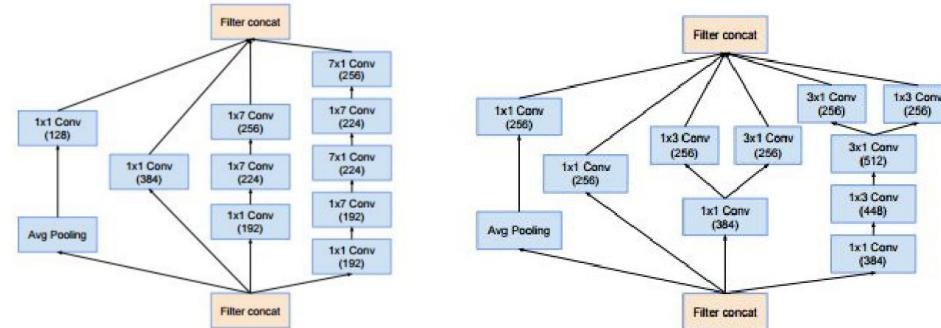
## **A Tutorial**

**Nikhil Naik**  
**nnaik@salesforce.com**

# Importance of Neural Architectures in Vision



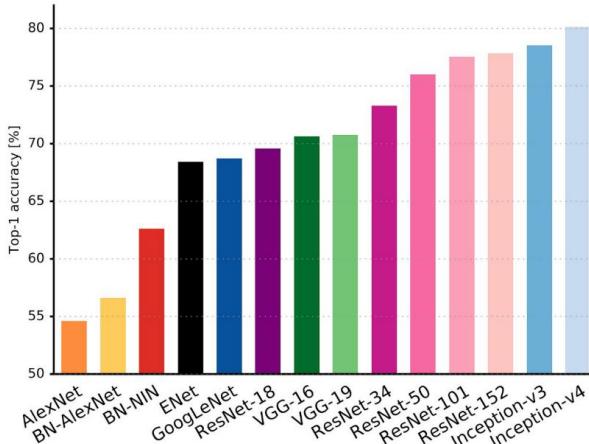
Canziani et al (2017)



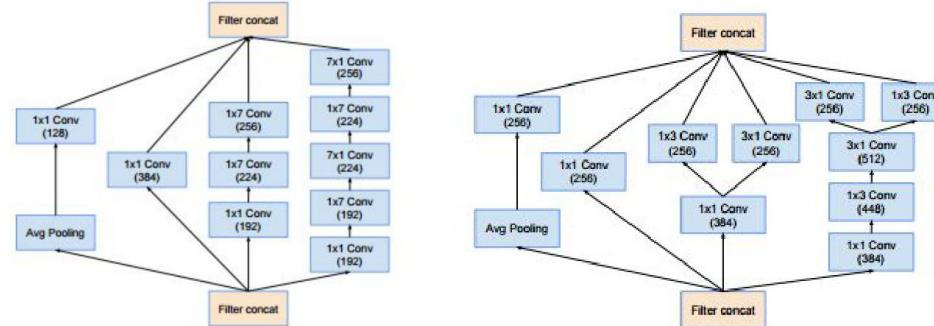
Complex hand-engineered layers from  
Inception-V4 (Szegedy et al., 2017)

**Design Innovations (2012 - Present):** Deeper networks, stacked modules, skip connections, squeeze-excitation block, ...

# Importance of Neural Architectures in Vision



Canziani et al (2017)



Complex hand-engineered layers from  
Inception-V4 (Szegedy et al., 2017)

Can we try and learn good architectures automatically?

# **Neural Architecture Search**

# Neural Architecture Search: Early Work

- **Neuroevolution:** Evolutionary algorithms (e.g., Miller et al., 89; Schaffer et al., 92; Stanley and Miikkulainen, 02; Verbancsics & Harguess, 13)
- Random search (e.g., Pinto et al., 09)
- Bayesian optimization for architecture and hyperparameter tuning (e.g., Snoek et al, 12; Bergstra et al., 13; Domhan et al., 15)

## **Renewed Interest in Neural Architecture Search (2017 -)**

# **NEURAL ARCHITECTURE SEARCH WITH REINFORCEMENT LEARNING**

**Barret Zoph,\* Quoc V. Le**  
Google Brain

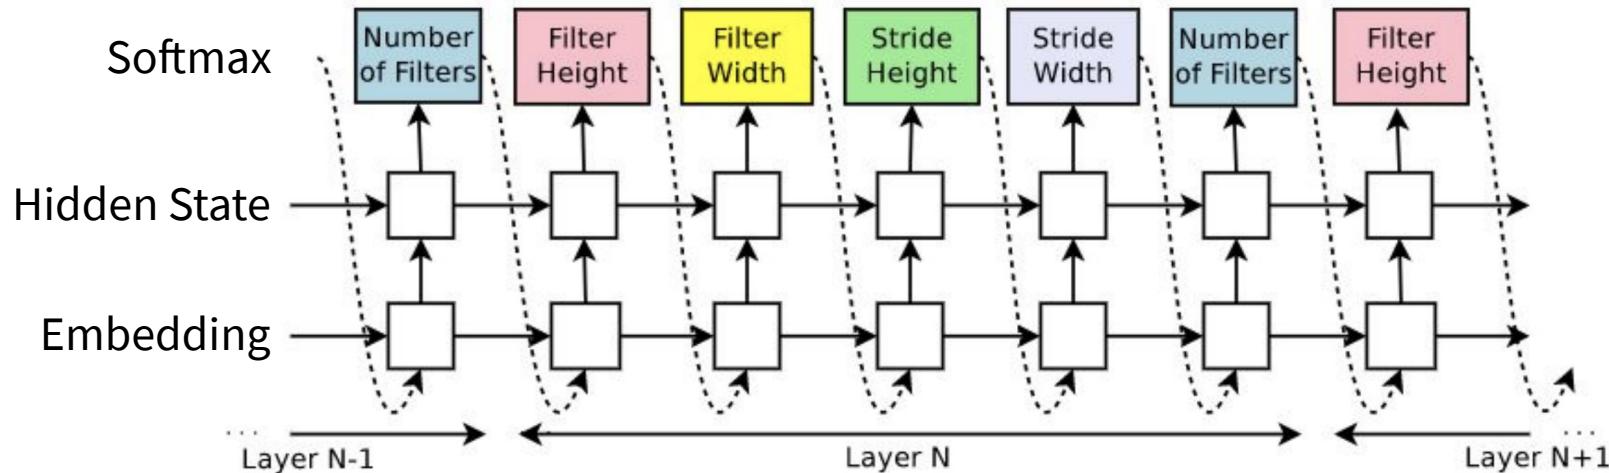
# **DESIGNING NEURAL NETWORK ARCHITECTURES USING REINFORCEMENT LEARNING**

**Bowen Baker, Otkrist Gupta, Nikhil Naik & Ramesh Raskar**  
Media Laboratory  
Massachusetts Institute of Technology

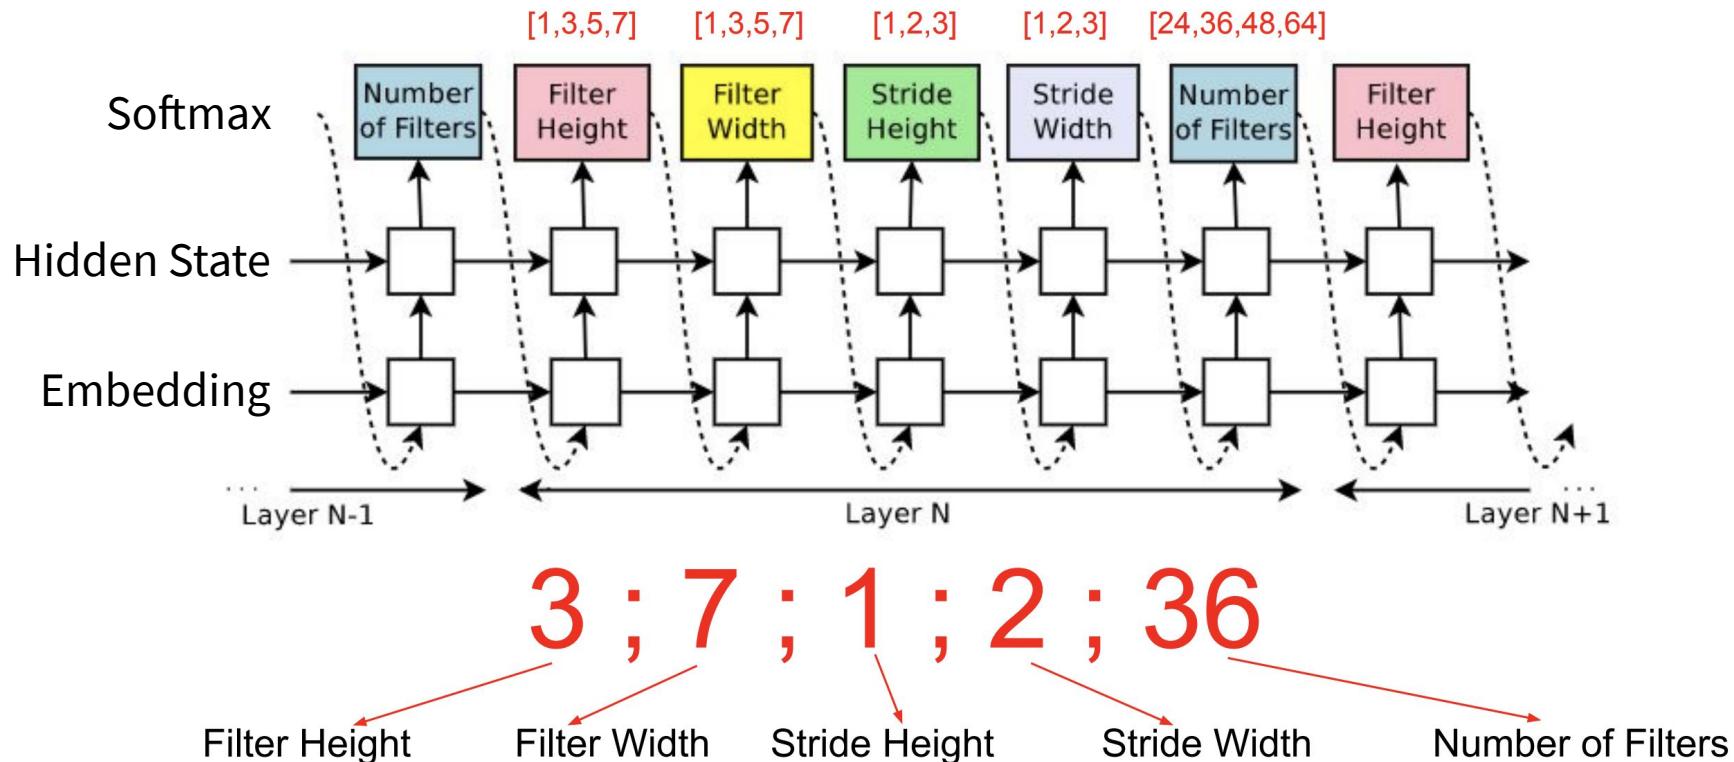
# Neural Architecture Search: Key Ideas

- Specify the structure and connectivity of a neural network by using a configuration string (e.g., [*“Filter Width: 5”*, *“Filter Height: 3”*, *“Num Filters: 24”*])
- Zoph and Le (2017): Use a RNN (“Controller”) to generate this string that specifies a neural network architecture
- Train this architecture (“Child Network”) to see how well it performs on a validation set
- Use reinforcement learning to update the parameters of the Controller model based on the accuracy of the child model

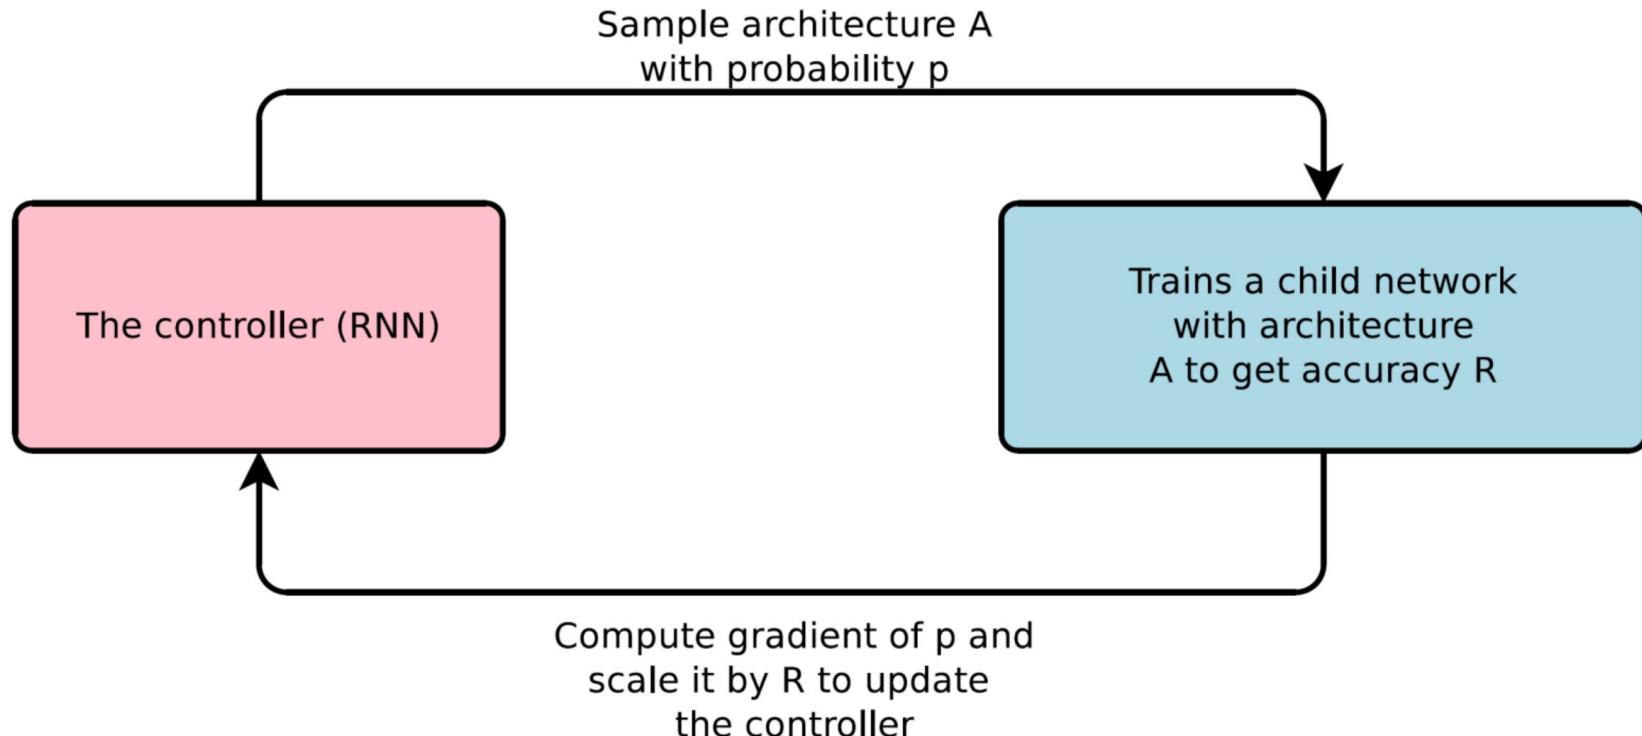
# Training with REINFORCE (Zoph and Le, 2017)



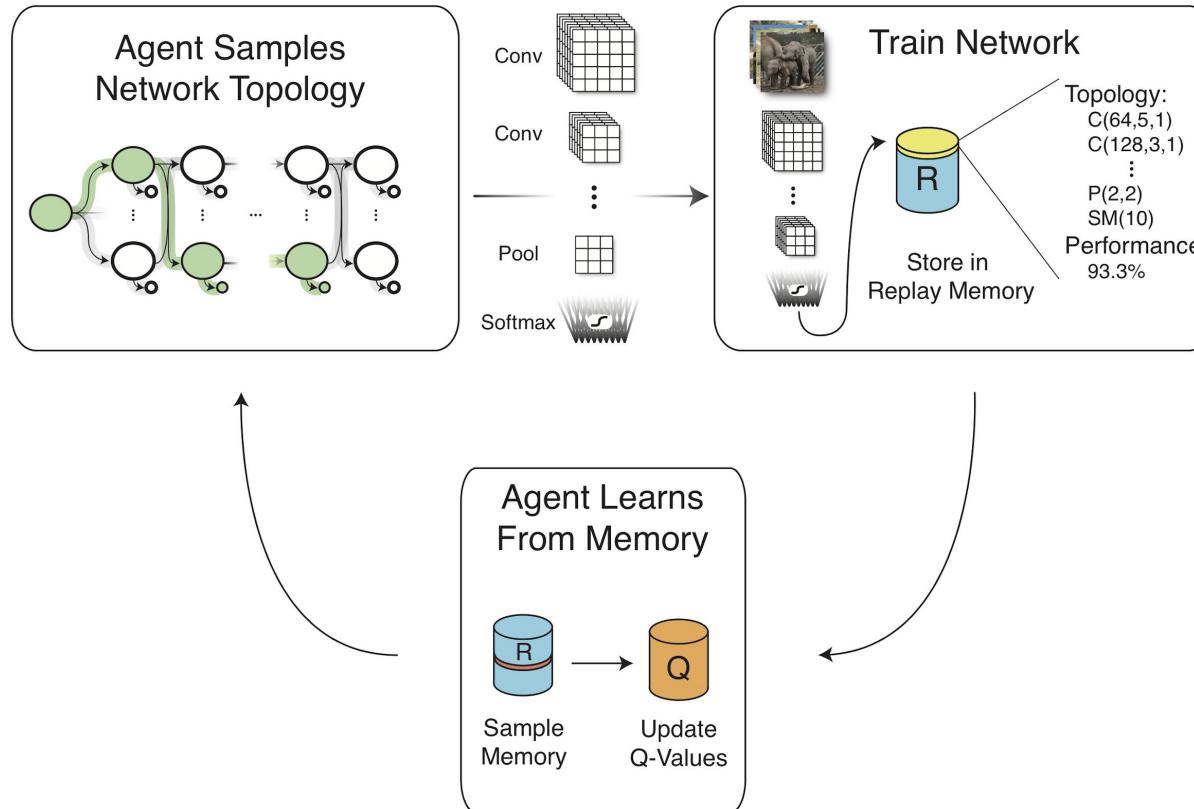
# Training with REINFORCE (Zoph and Le, 2017)



# Training with REINFORCE (Zoph and Le, 2017)

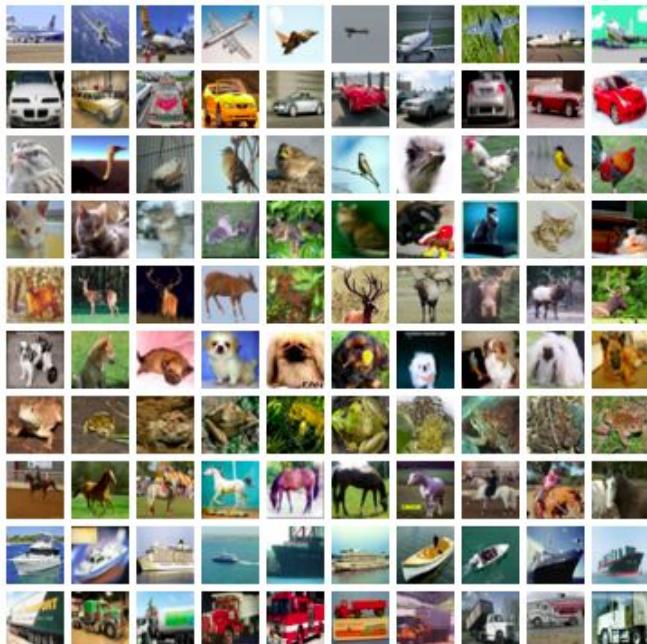


# Q-Learning with Experience Replay (Baker et al., 2017)



# Results on CIFAR-10

**Zoph and Le (2017):** 12,800 architectures trained using 22,400 GPUdays  
**Baker et al. (2017):** 2,700 architectures trained using 100 GPUdays

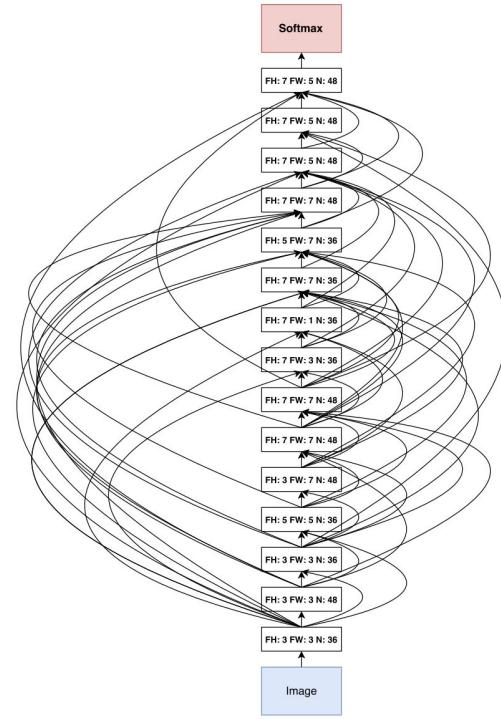


50K Training Samples  
10K Testing Samples  
10 classes

# Results on CIFAR-10

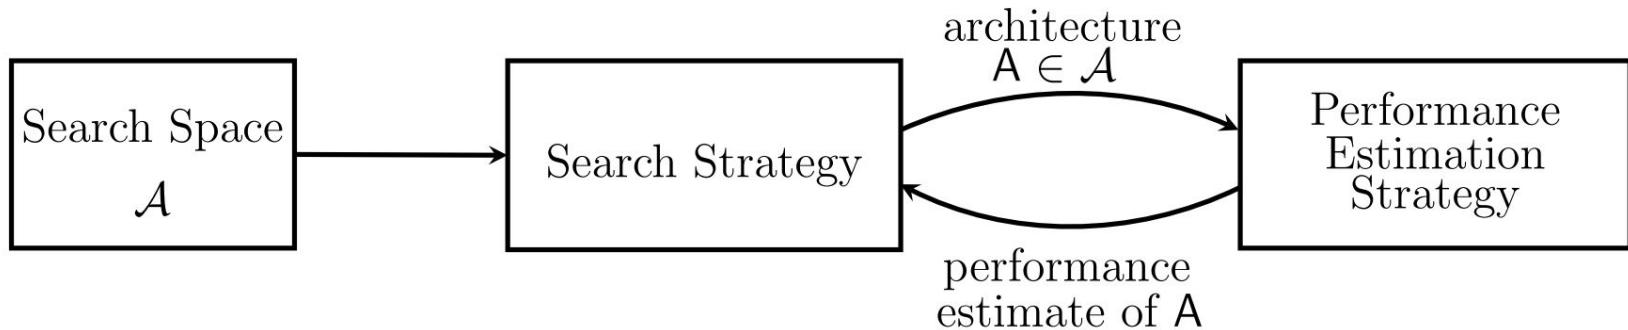
Model	Depth	Parameters	Error rate (%)
Wide ResNet (Zagoruyko & Komodakis, 2016)	16	11.0M	4.81
	28	36.5M	4.17
ResNet (pre-activation) (He et al., 2016b)	164	1.7M	5.46
	1001	10.2M	4.62
DenseNet ( $L = 40, k = 12$ ) (Huang et al., 2016a)	40	1.0M	5.24
DenseNet ( $L = 100, k = 12$ ) (Huang et al., 2016a)	100	7.0M	4.10
DenseNet ( $L = 100, k = 24$ ) (Huang et al., 2016a)	100	27.2M	3.74
DenseNet-BC ( $L = 100, k = 40$ ) (Huang et al., 2016b)	190	25.6M	3.46
Neural Architecture Search v1 no stride or pooling	15	4.2M	5.50
Neural Architecture Search v2 predicting strides	20	2.5M	6.01
Neural Architecture Search v3 max pooling	39	7.1M	4.47
Neural Architecture Search v3 max pooling + more filters	39	37.4M	3.65
MetaQNN (Baker et al. 2017) no skip connections	12	11.8M	6.92

Comparable accuracy to best  
human-designed models ~2017

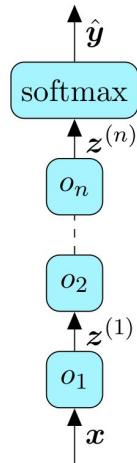


Best NAS Architecture on CIFAR-10  
Zoph and Le (2017)

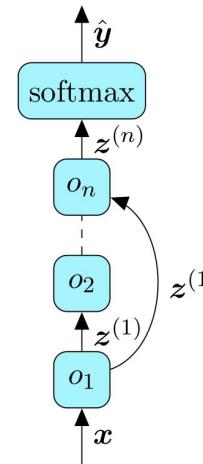
# Neural Architecture Search: General Problem Setup



# NAS: Global Search Space



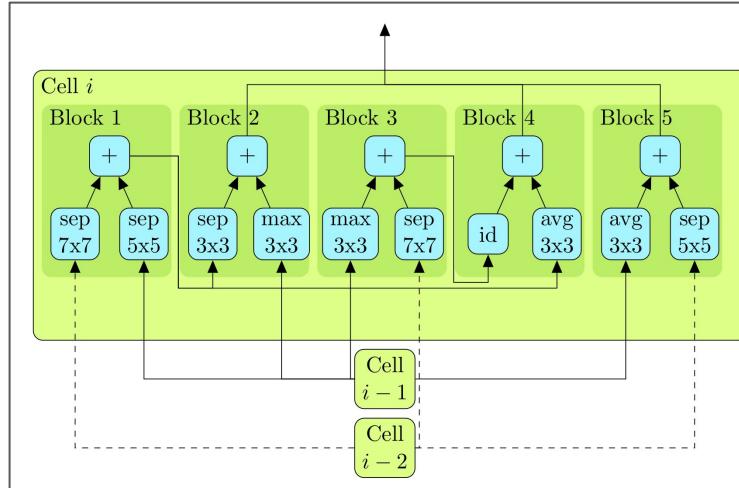
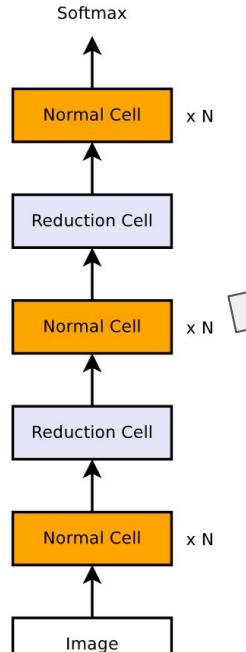
MetaQNN  
(Baker et al. 2017)



NAS  
(Zoph and Le, 2017)

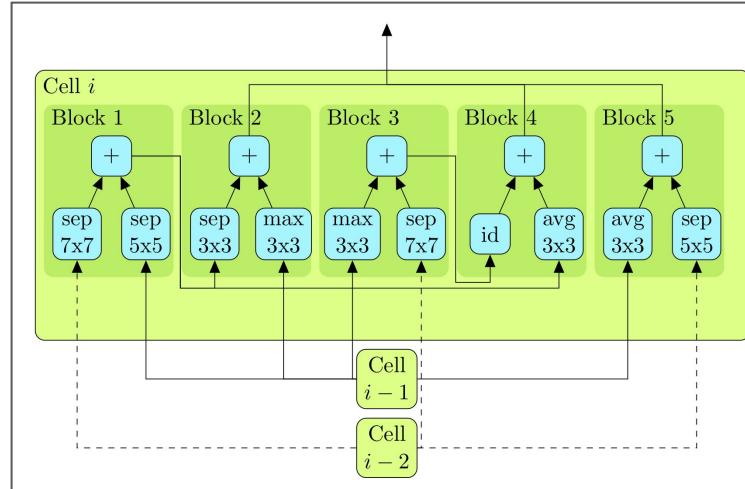
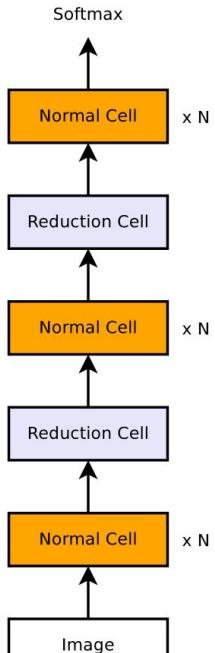
**Global search space is rigid, impractical to scale and transfer**

# NAS: Cell-based Search Space

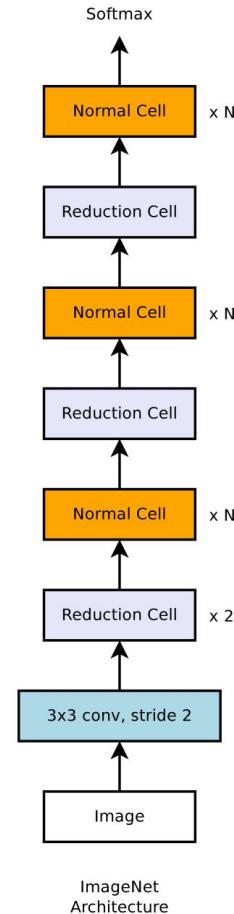


**NASNet**  
Zoph et al. (2018)

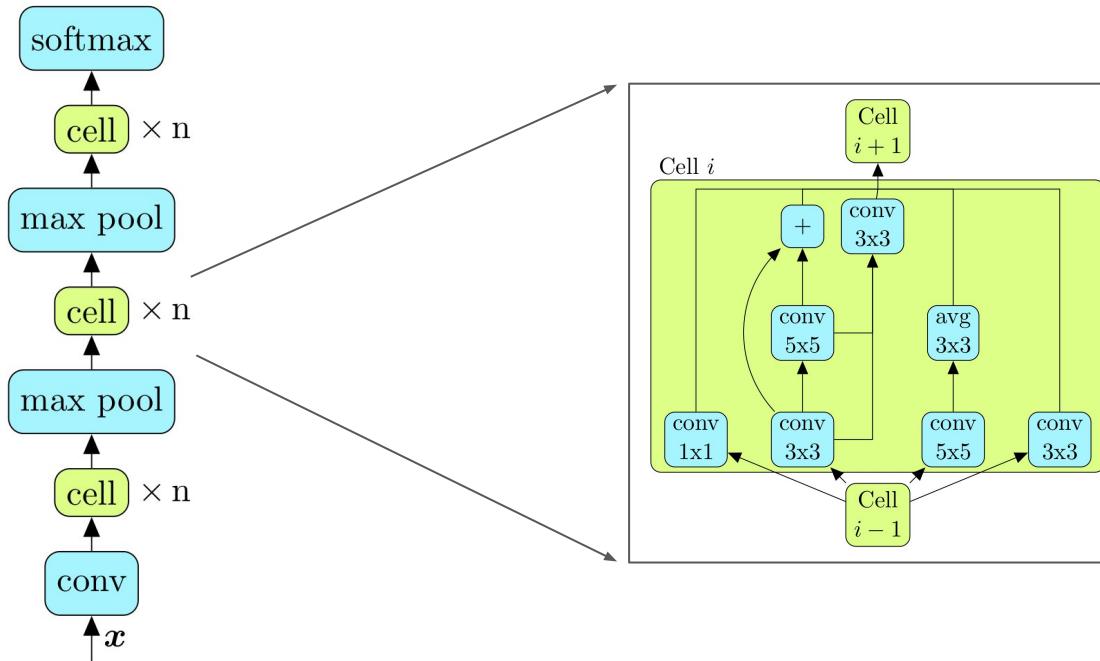
# NAS: Cell-based Search Space



**NASNet**  
**Zoph et al. (2018)**



# NAS: Cell-based Search Space

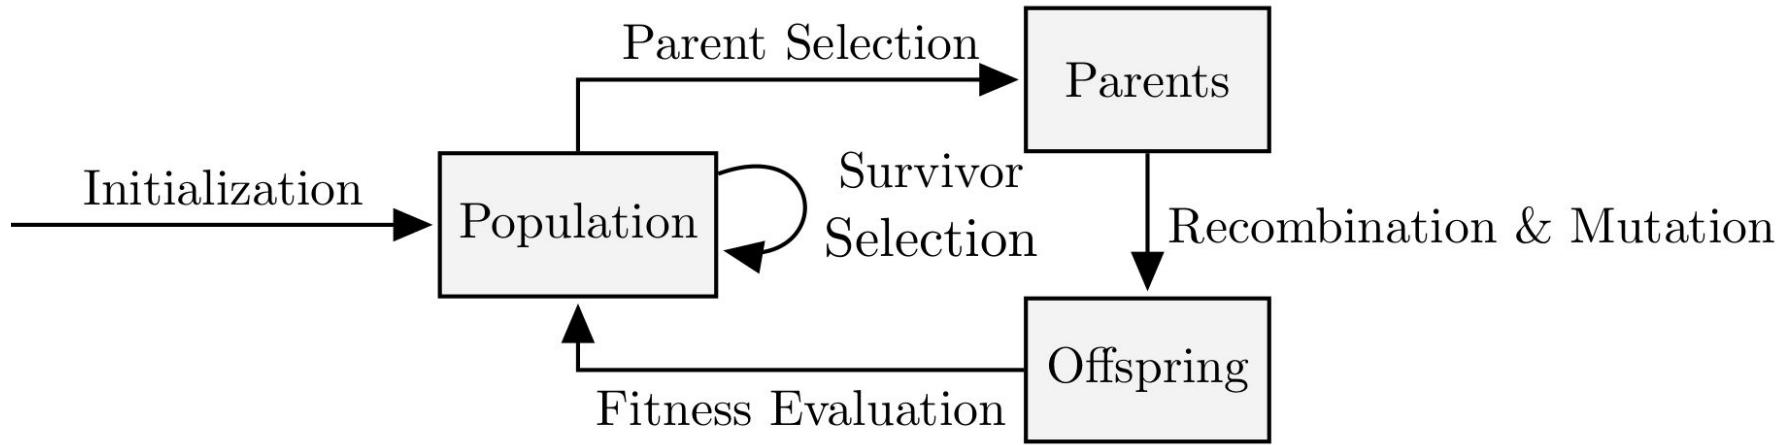


**BlockQNN**  
Zong et al. (2018)

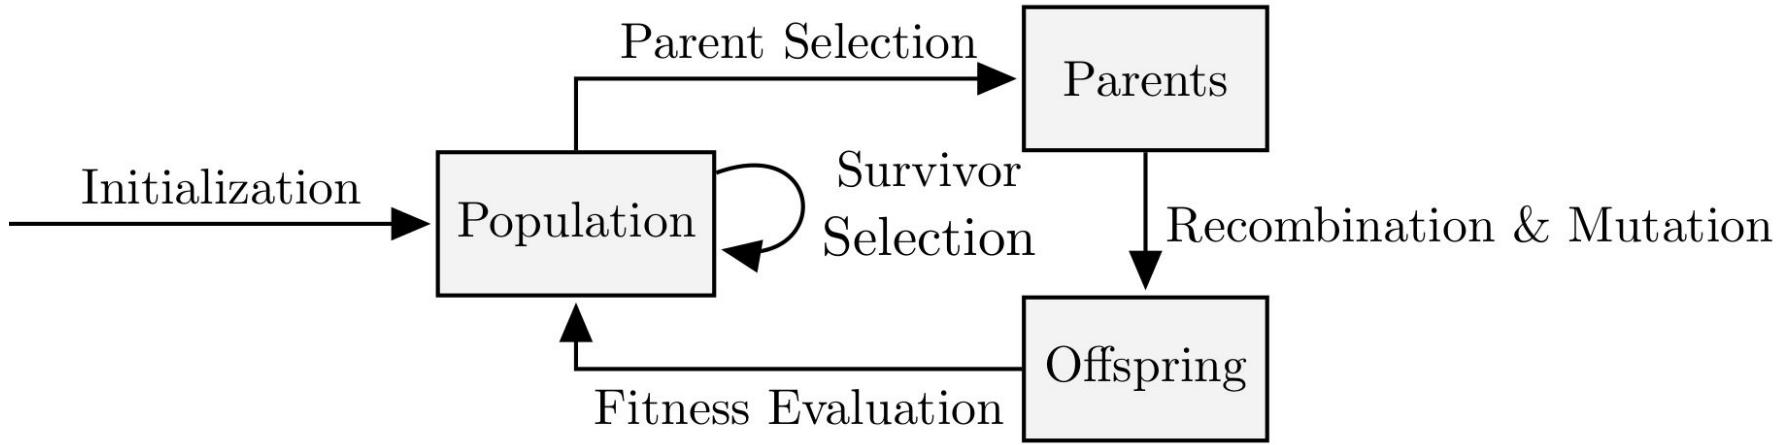
# Transferring Architectures from CIFAR-10 to ImageNet

Search Method	Search Space	Search Strategy	Search Cost (GPU-days)	CIFAR-10 Error	ImageNet Top-1 Error
MetaQNN (Baker et al, 2017)	Global	Q-learning	100	6.92	-
NAS (Zoph and Le, 2017)	Global	REINFORCE	22,400	3.65	-
BlockQNN (Zong et al, 2018)	Cell-based	Q-learning	96	3.54	<b>22.6</b>
NASNet (Zoph et al, 2018)	Cell-based	PPO	2,000	3.41	<b>17.3</b>

# Search Strategy: Evolution

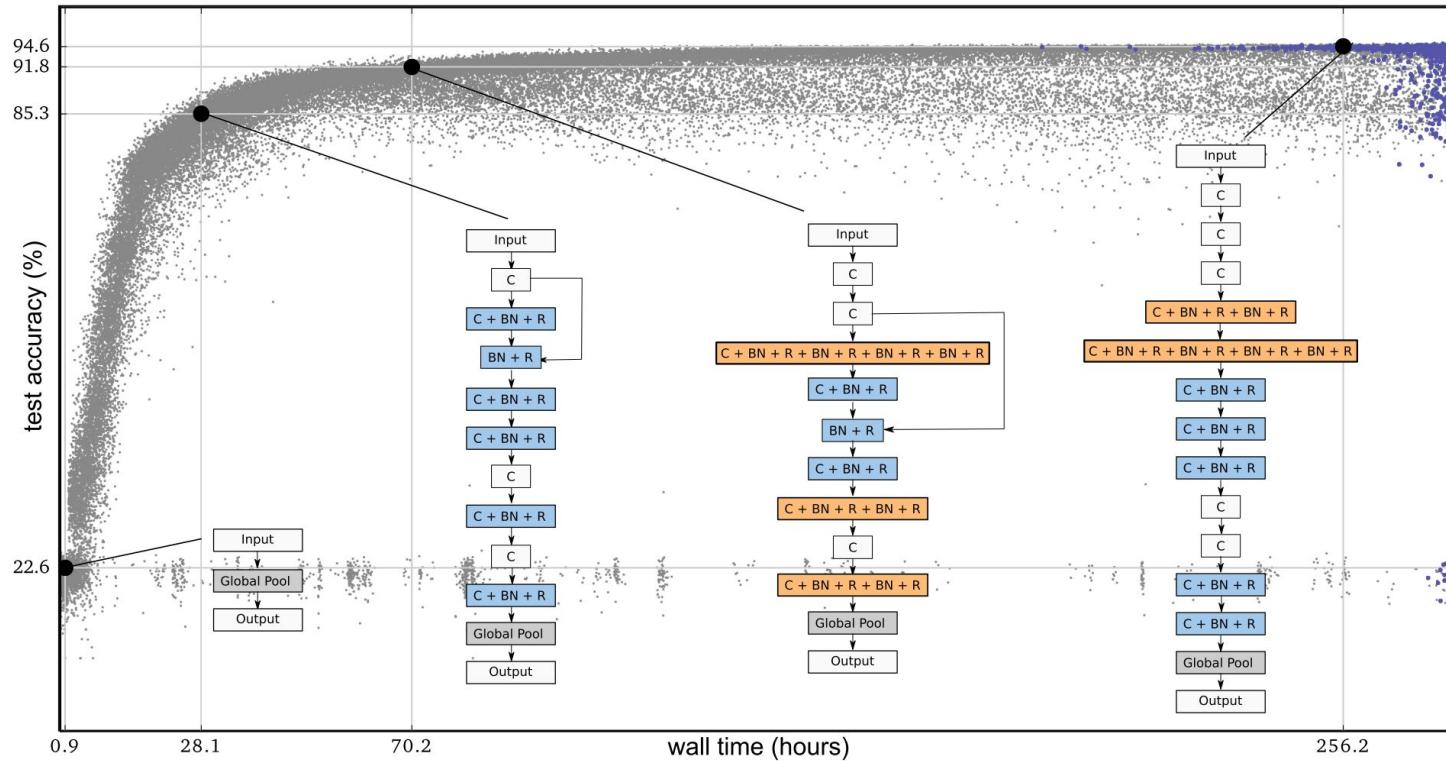


# Search Strategy: Evolution



**Typical Policies for Parent Selection and Offspring Retention**  
Tournament, Fitness-proportionate selection, Youngest

# Evolution over Global Search Space on CIFAR-10



Real et al. (2017)

# Evolution: Results on CIFAR-10

Method	Parent Sel.	Survivor Sel.	Error (%)
Real et al. (2017)	tournament	tournament	5.40
Xie and Yuille (2017)	all	fitness prop.	5.39
Suganuma et al. (2017)	n/a	elitist	5.98
Liu et al. (2018b)	tournament	all	3.75
Real et al. (2019)	tournament	youngest	3.34
Elsken et al. (2018)	n/a	elitist	5.2
Wistuba (2018)	tournament	all	3.57

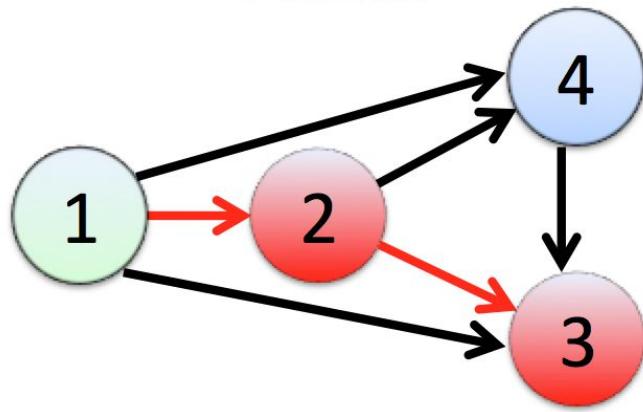
# Computational Cost of NAS on CIFAR-10

Reference	Error (%)	Params (Millions)	GPU Days	
Baker et al. (2017)	6.92	11.18	100	Reinforcement Learning
Zoph and Le (2017)	3.65	37.4	22,400	
Cai et al. (2018a)	4.23	23.4	10	
Zoph et al. (2018)	3.41	3.3	2,000	
Zoph et al. (2018) + Cutout	2.65	3.3	2,000	
Zhong et al. (2018)	3.54	39.8	96	
Cai et al. (2018b)	2.99	5.7	200	
Cai et al. (2018b) + Cutout	2.49	5.7	200	
Real et al. (2017)	5.40	5.4	2,600	Evolution
Xie and Yuille (2017)	5.39	N/A	17	
Suganuma et al. (2017)	5.98	1.7	14.9	
Liu et al. (2018b)	3.75	15.7	300	
Real et al. (2019)	3.34	3.2	3,150	

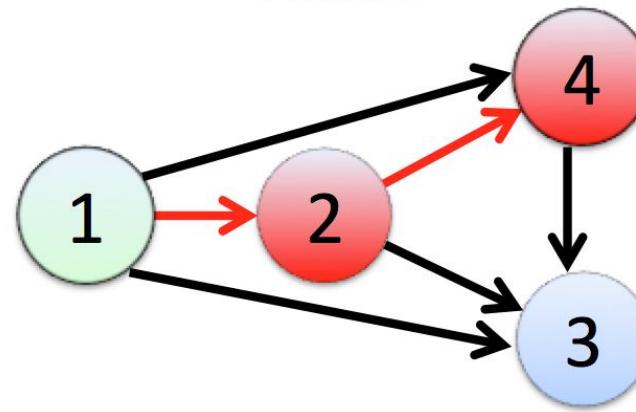
**Designing competitive networks can take hundreds of GPU-days!**  
**How to make neural architecture search more efficient?**

# How To Make NAS More Efficient?

Path A



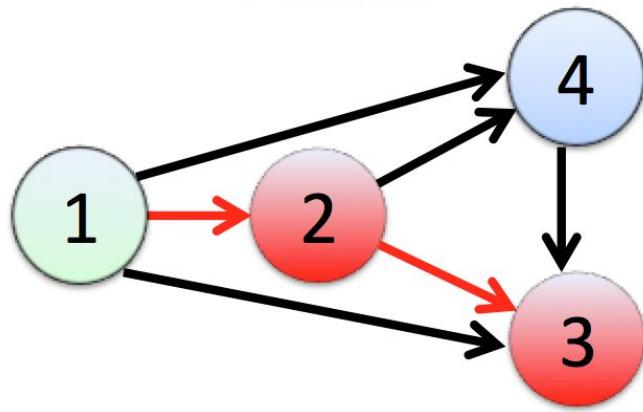
Path B



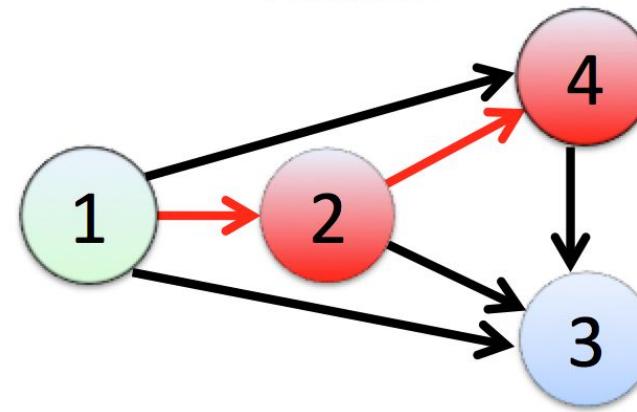
- Currently, models defined by **path A** and **path B** are trained independently
- Instead, treat all model trajectories as sub-graphs of a single directed acyclic graph
- Use a search strategy (e.g., RL, Evolution) to choose sub-graphs. Proposed in ENAS (Pham et al, 2018)

# Efficient NAS using Weight Sharing

Path A

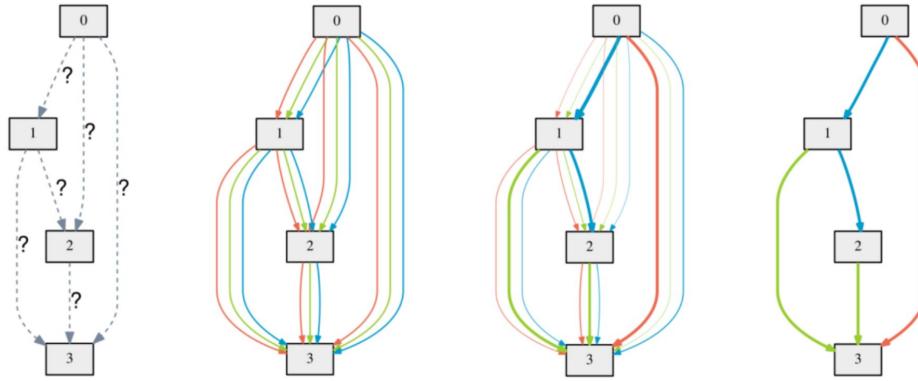


Path B



- Currently, models defined by **path A** and **path B** are trained independently
- Instead, treat all model trajectories as sub-graphs of a single directed acyclic graph
- Use a search strategy (e.g., RL, Evolution) to choose sub-graphs. Proposed in ENAS (Pham et al, 2018)

# Gradient-based NAS with Weight Sharing



**DARTS (Liu et al., 2018)**

Create a mixed operation  $\bar{o}^{(i,j)}$  parametrized by  $\alpha^{(i,j)}$  for each edge  $(i, j)$

**while** *not converged* **do**

1. Update architecture  $\alpha$  by descending  $\nabla_\alpha \mathcal{L}_{val}(w - \xi \nabla_w \mathcal{L}_{train}(w, \alpha), \alpha)$   
( $\xi = 0$  if using first-order approximation)
2. Update weights  $w$  by descending  $\nabla_w \mathcal{L}_{train}(w, \alpha)$

**Derive the final architecture based on the learned  $\alpha$ .**

Find encoding weights  $a^*$  that

$$\min_{\alpha} \mathcal{L}_{val}(w^*(\alpha), \alpha)$$

$$\text{s.t. } w^*(\alpha) = \operatorname{argmin}_w \mathcal{L}_{train}(w, \alpha)$$

Learn design of normal and reduction cells

# Efficient NAS with Weight Sharing: Results on CIFAR-10

Reference	Error (%)	Params (Millions)	GPU Days
Pham et al. (2018)	3.54	4.6	0.5
Pham et al. (2018) + Cutout	2.89	4.6	0.5
Bender et al. (2018)	4.00	5.0	N/A
Casale et al. (2019) + Cutout	2.81	3.7	1
Liu et al. (2018c) + Cutout	2.76	3.3	4
Xie et al. (2019b) + Cutout	2.85	2.8	1.5
Cai et al. (2019) + Cutout	2.08	5.7	8.33
Brock et al. (2018)	4.03	16.0	3
Zhang et al. (2019)	4.30	5.1	0.4

## Limitations:

- Restrict the search space to the subgraphs of the supergraph
- Can bias the search towards certain regions of the search space
- In practice, the need to hold entire supergraph in GPU memory restricts search space size

# How Well Do Automatically Designed Architectures Transfer To Other Datasets?

# CIFAR-10 → ImageNet Architecture Transfer

Model	Params	×+	1/5-Acc (%)
Inception V3	23.8M	5.72B	78.8 / 94.4
Xception	22.8M	8.37B	79.0 / 94.5
Inception ResNet V2	55.8M	13.2B	80.4 / 95.3
ResNeXt-101 (64x4d)	83.6M	31.5B	80.9 / 95.6
PolyNet	92.0M	34.7B	81.3 / 95.8
Dual-Path-Net-131	79.5M	32.0B	81.5 / 95.8
Squeeze-Excite-Net	145.8M	42.3B	82.7 / 96.2
GeNet-2*	156M	—	72.1 / 90.4
Block-QNN-B (N=3)*	—	—	75.7 / 92.6
Hierarchical (2, 64)*	64M	—	79.7 / 94.8
PNASNet-5 (4, 216)	86.1M	25.0B	82.9 / 96.1
NASNet-A (6, 168)	88.9M	23.8B	82.7 / 96.2
AmoebaNet-B (6, 190)*	84.0M	22.3B	82.3 / 96.1
AmoebaNet-A (6, 190)*	86.7M	23.1B	82.8 / 96.1
AmoebaNet-A (6, 204)*	99.6M	26.2B	82.8 / 96.2
AmoebaNet-C (6, 228)*	155.3M	41.1B	<b>83.1 / 96.3</b>

# Architecture Transferability to Other Datasets and Tasks

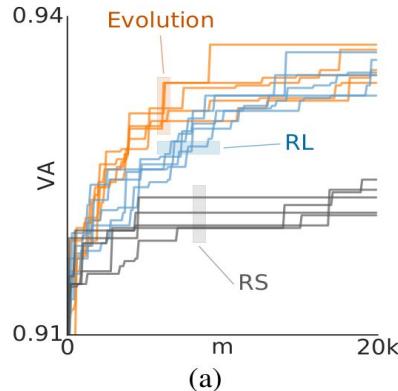
Dataset	Acc.	Network	Acc.	Best network
Food-101	<b>90.0</b>	Deep layer aggregation [40]	<b>90.1</b>	NASNet-A Large, fine-tuned ✓
CIFAR-10	97.9	AmoebaNet [41]	<b>98.4<sup>a</sup></b>	NASNet-A Large, fine-tuned ✓
CIFAR-100	87.8	ShakeDrop [42]	<b>88.2<sup>a</sup></b>	NASNet-A Large, fine-tuned ✓
Birdsnap	<b>80.2<sup>b</sup></b>	Mask-CNN [43]	78.5	NASNet-A Large, fine-tuned ✓
SUN397	63.2	Places-pretrained VGG [44]	<b>66.5</b>	NASNet-A Large, fine-tuned ✓
Stanford Cars	<b>94.1</b>	Deep layer aggregation [40]	93.0	Inception v4, random init
FGVC Aircraft	<b>92.9<sup>b</sup></b>	Deep layer aggregation [40]	89.4	Inception v3, fine-tuned
VOC 2007 Cls.	<b>89.7</b>	VGG [9]	88.4	NASNet-A Large, fine-tuned ✓
DTD	75.5	FC+FV-CNN+D-SIFT [45]	<b>76.7</b>	Inception-ResNet v2, fine-tuned
Oxford-IIIT Pets	93.8	Object-part attention [46]	<b>94.3</b>	NASNet-A Large, fine-tuned ✓
Caltech-101	93.4	Spatial pyramid pooling [47]	<b>95.0</b>	NASNet-A Large, fine-tuned ✓
Oxford 102 Flowers	97.1	Object-part attention [46]	<b>97.7</b>	NASNet-A Large, fine-tuned ✓

NASNet-A: SOTA on **8/13 commonly used classification benchmarks**

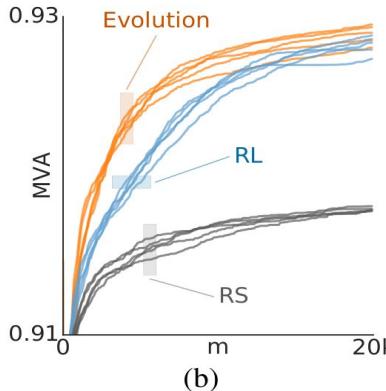
NAS architectures also transfer for object detection and semantic segmentation

# **Are Neural Architecture Search Strategies better than Random Search?**

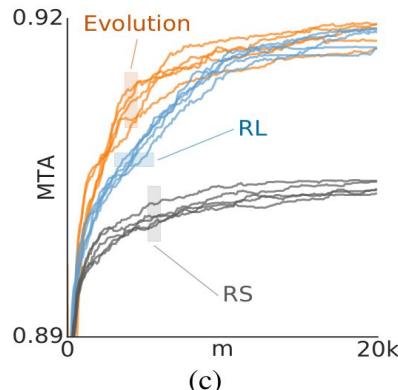
# Are Intelligent Search Strategies Better Than Random Search?



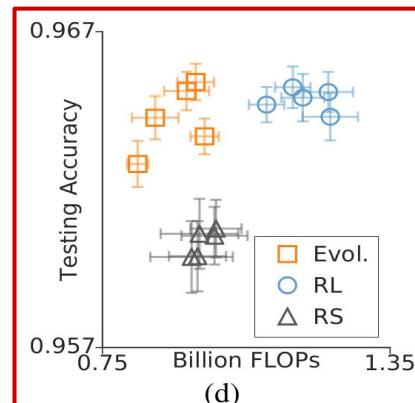
(a)



(b)



(c)



**Real et al. (2018)**  
The difference in accuracy  
between best models  
found by random search,  
RL, and Evolution is  
**less than 1%** on CIFAR-10

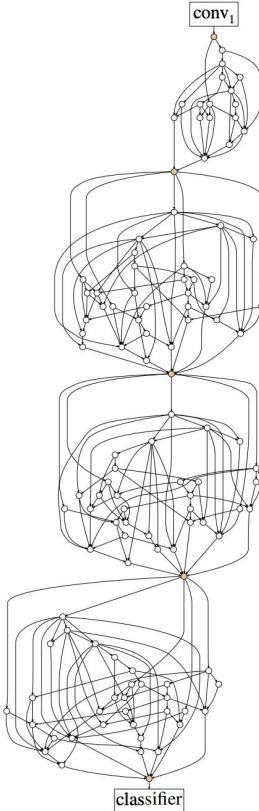
# Are Intelligent Search Strategies Better Than Random Search?

Architecture	Source	Test Error		Params (M)
		Best	Average	
NASNet-A <sup>#*</sup>	[52]	N/A	2.65	3.3
AmoebaNet-B <sup>*</sup>	[43]	N/A	$2.55 \pm 0.05$	2.8
ProxylessNAS <sup>†</sup>	[7]	2.08	N/A	5.7
GHN <sup>#†</sup>	[50]	N/A	$2.84 \pm 0.07$	5.7
SNAS <sup>†</sup>	[47]	N/A	$2.85 \pm 0.02$	2.8
ENAS <sup>†</sup>	[41]	2.89	N/A	4.6
ENAS	[34]	2.91	N/A	4.2
Random search baseline	[34]	N/A	$3.29 \pm 0.15$	3.2
DARTS (first order)	[34]	N/A	$3.00 \pm 0.14$	3.3
DARTS (second order)	[34]	N/A	$2.76 \pm 0.09$	3.3
DARTS (second order) <sup>‡</sup>	Ours	2.62	$2.78 \pm 0.12$	3.3
ASHA baseline	Ours	2.85	$3.03 \pm 0.13$	2.2
Random search WS <sup>‡</sup>	Ours	2.71	$2.85 \pm 0.08$	4.3

**Li and Talwalkar (2019)**

Random search baseline finds a model with 2.71% error on CIFAR-10, comparable to best NAS methods based on RL, Evolution, Gradient Descent

# Are Intelligent Search Strategies Better Than Random Search?

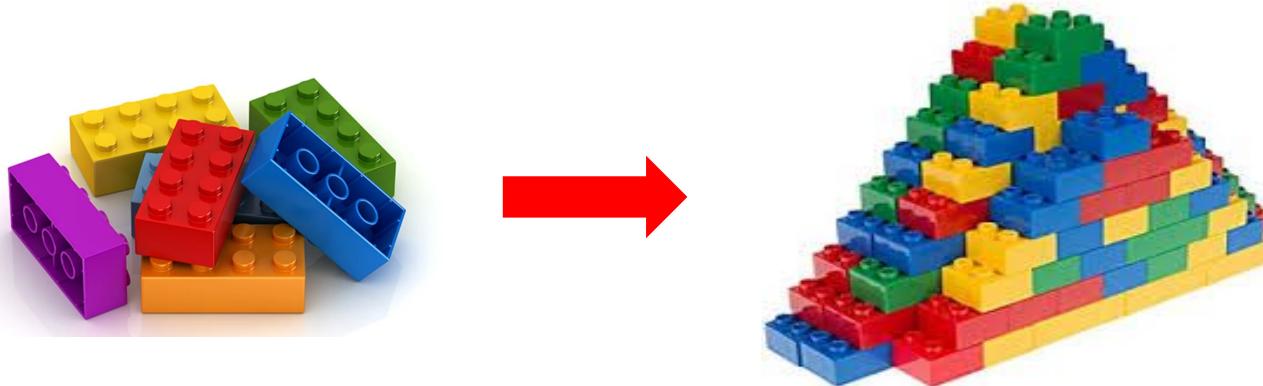


**Xie et al. (2019)**

A network consisting of multiple randomly wired “cells” is only **1.3% less accurate** than a similar capacity NAS models on **ImageNet**

network	test size	epochs	top-1 acc.	top-5 acc.	FLOPs (B)	params (M)
NASNet-A [56]	$331^2$	>250	82.7	96.2	23.8	88.9
Amoeba-B [34]	$331^2$	>250	82.3	96.1	22.3	84.0
Amoeba-A [34]	$331^2$	>250	82.8	96.1	23.1	86.7
PNASNet-5 [26]	$331^2$	>250	82.9	96.2	25.0	86.1
<b>RandWire-WS</b>	$320^2$	100	$81.6 \pm 0.13$	$95.6 \pm 0.07$	$16.0 \pm 0.36$	$61.5 \pm 1.32$

# Are Intelligent Search Strategies Better Than Random Search?



- Random search is a competitive baseline
- How do you design more flexible search spaces and sample-efficient search methods?
- Can intelligent search methods help discover new design motifs and basic building blocks?

# Designing Efficient Architectures

# Designing Efficient Architectures

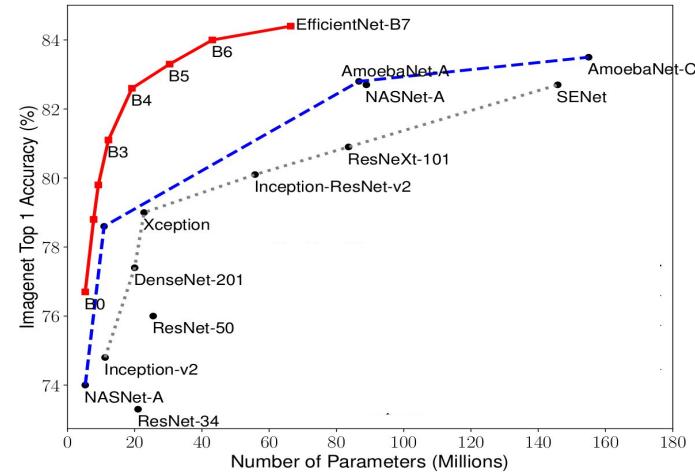
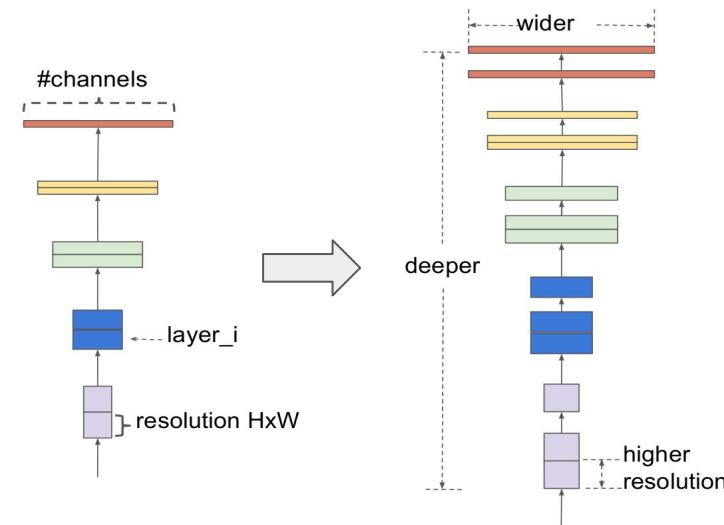
- **Design high-performing models given a budget on memory and inference time**
- Many applications in resource constrained-settings: Mobile devices, autonomous vehicles, robotics, production models
- Popular approach: Tweak existing architectures with weight pruning (e.g. LeCun et al. '89, Huang et al. '18) or weight quantization (e.g., Binarized neural nets '16, XNORnet '16)
- Neural Architecture Search is well-suited for this application

# Designing Efficient Architectures

Introduce constraints like memory and inference time in addition to accuracy

NAS Method	References
Constrained Optimization	Tan et al. (2018), Cai et al. (2018), Hsu et al. (2018)
Multi-objective Optimization	Kim et al. (2017), Cai et al. (2019), Lu et al. (2018), Dong et al. (2018), Elsken et al. (2019), Tan and Le (2019)
Automated Pruning	He et al. (2019)

# Designing Efficient Architectures from Scratch

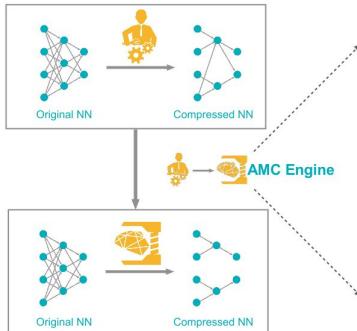


## EfficientNet (Tan and Le, 2019)

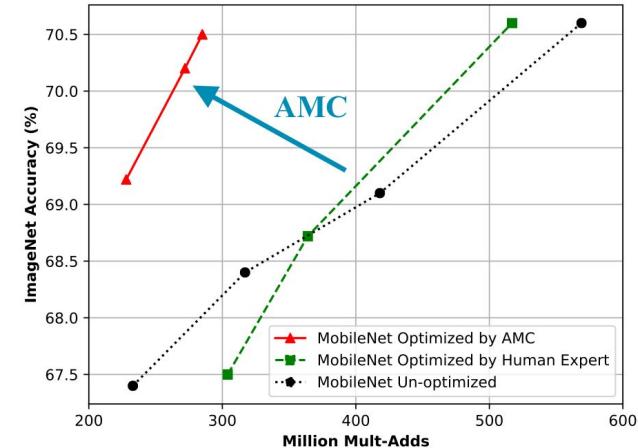
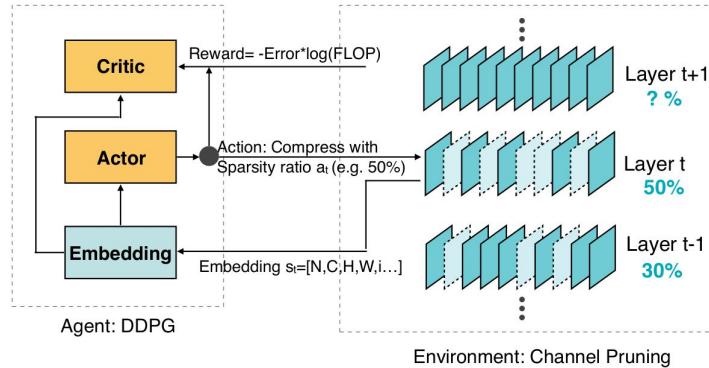
RL-based architecture search for a feed-forward block + scaling up using grid search  
**80% Top-1 Accuracy on ImageNet with 5x fewer parameters and 13x fewer FLOPS**  
than best human-designed model

# Designing Efficient Architectures: Auto-Pruning

Model Compression by Human:  
Labor Consuming, Sub-optimal



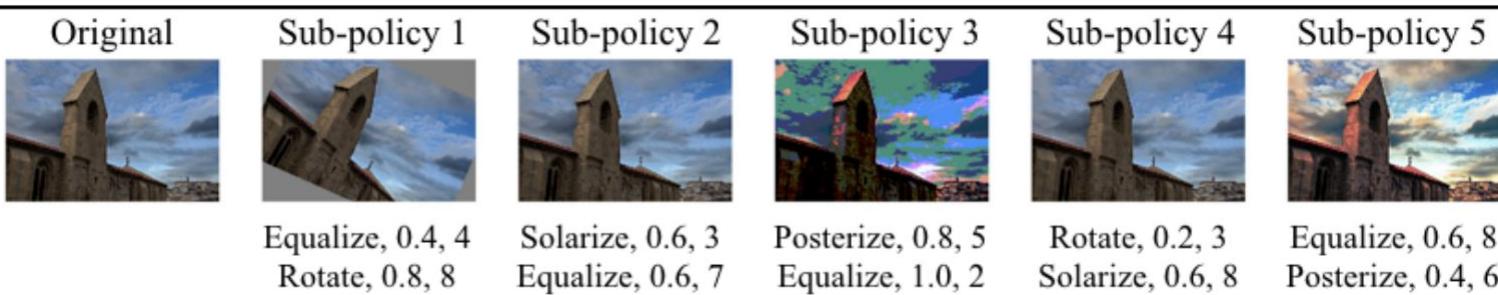
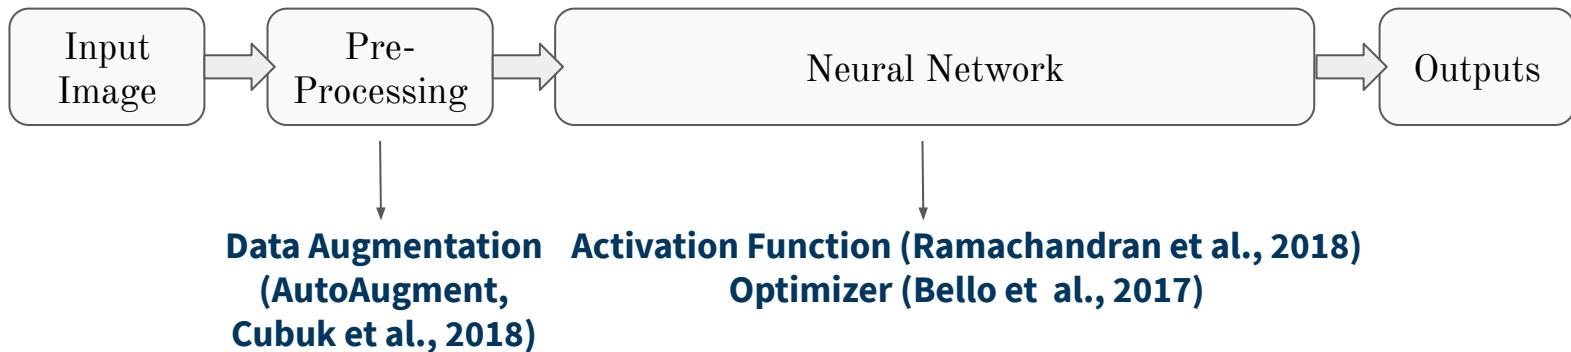
Model Compression by AI:  
Automated, Higher Compression Rate, Faster



AMC: AutoML for Model Compression (He et al., 2019)  
Using Deep Deterministic Policy Gradients to learn pruning ratio for each layer

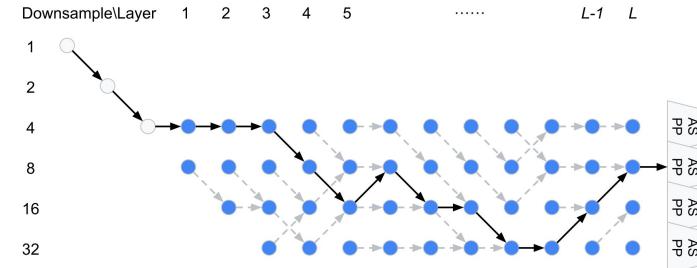
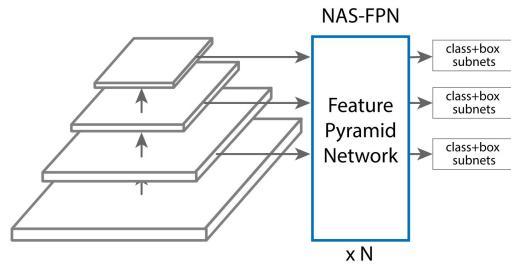
# Automating the Entire Deep Learning Stack

# Automating the Deep Learning Stack



Model	Baseline	Inception Pre-processing [14]	AutoAugment
ResNet-50 [15]	24.70 / 7.80	23.69 / 6.92	22.37 / 6.18
ResNet-200 [15]	-	21.52 / 5.85	20.00 / 4.99

# Neural Architecture Search: Beyond Image Classification

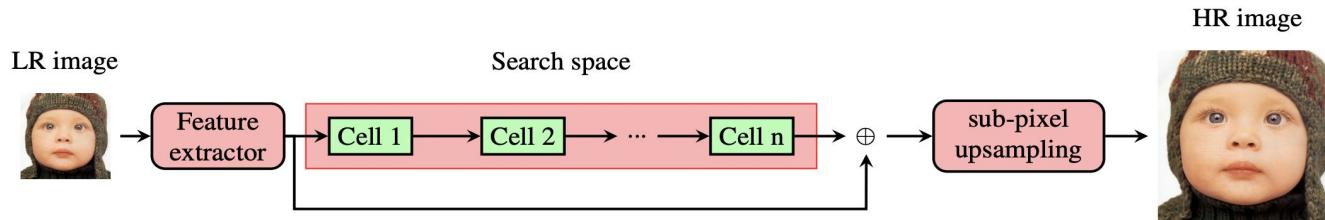


## Object Detection

E.g. NAS-FPN (Ghaisi et al., 2019)

## Semantic Segmentation

E.g. Auto-DeepLab (Liu et al., 2019)



**Image Super-Resolution** e.g. MoreMNAS(Chu et al., 2019)

# Summary

- **Neural architecture search** can generate state-of-the-art architectures for computer vision tasks given a well-defined search space
- Develop new optimizers, activation functions, data augmentation strategies
- Architecture search can be especially useful for designing efficient architectures for resource-constrained settings

# Future Directions

- Search efficiency
- Moving towards less constrained search spaces
- Designing efficient architectures: Automated scaling, pruning and quantization
- Joint optimization of all components in the deep learning pipeline (data augmentation, architectures, activation functions, training algorithms)
- Designing architectures for multimodal problems, e.g., vision and language

## Useful Survey Papers and Datasets

Elsken, Thomas, Jan Hendrik Metzen, and Frank Hutter. "Neural architecture search: A survey." *JMLR* (2019)

Wistuba, Martin, Ambrish Rawat, and Tejaswini Pedapati. "A Survey on Neural Architecture Search." *arXiv preprint arXiv:1905.01392* (2019)

Ying, Chris, Aaron Klein, Esteban Real, Eric Christiansen, Kevin Murphy, and Frank Hutter. "NAS-Bench-101: Towards reproducible neural architecture search." *ICML* (2019)

# thank you

